


ECE (0x440EC000)

Introduction to SystemVerilog for Design and Verification

# Original Floating Point Adder Design

```
0100 0100 0000 1110 1100 0000 0000 0000
```

Chris Kane-Pardy & Violet Monti

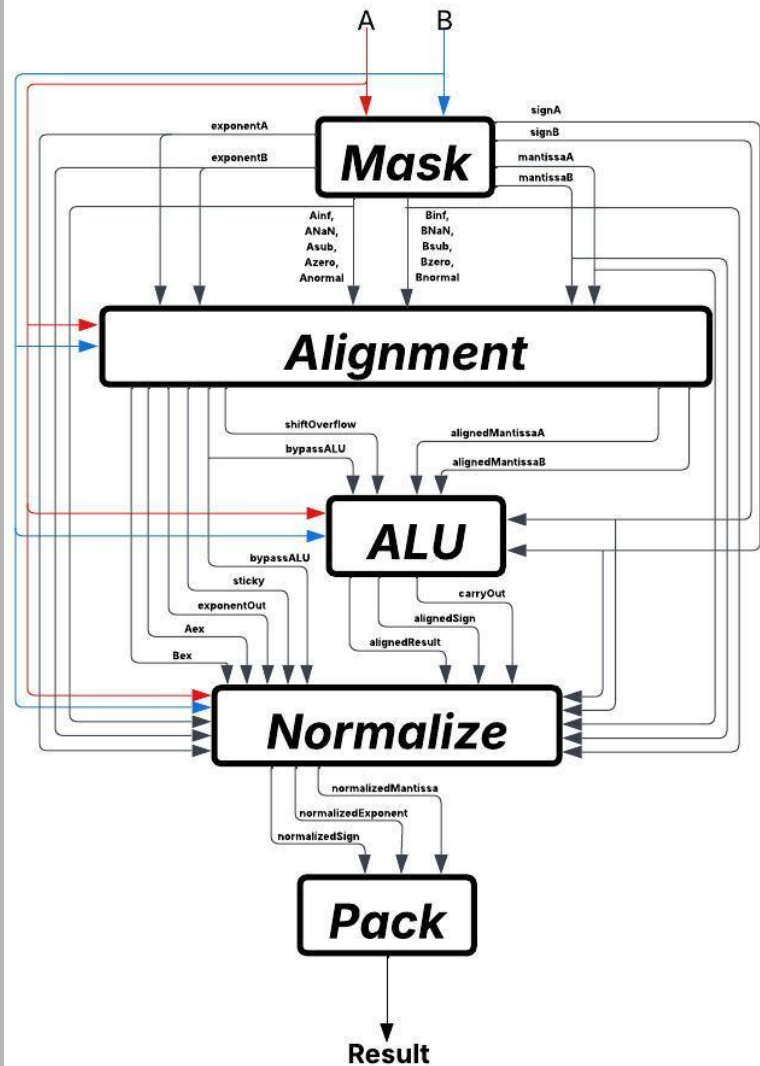


## Overview

- Designed 5 submodules for Easy incremental Testing
- Easy support for the addition of different rounding modes (package)
- Decided not to parameterize due to lack of 16-bit floating point data type for accurate comparison
- Special handling for non-normal numbers
- IEEE-754 compliant
- Fully asynchronous to allow use in pipelines

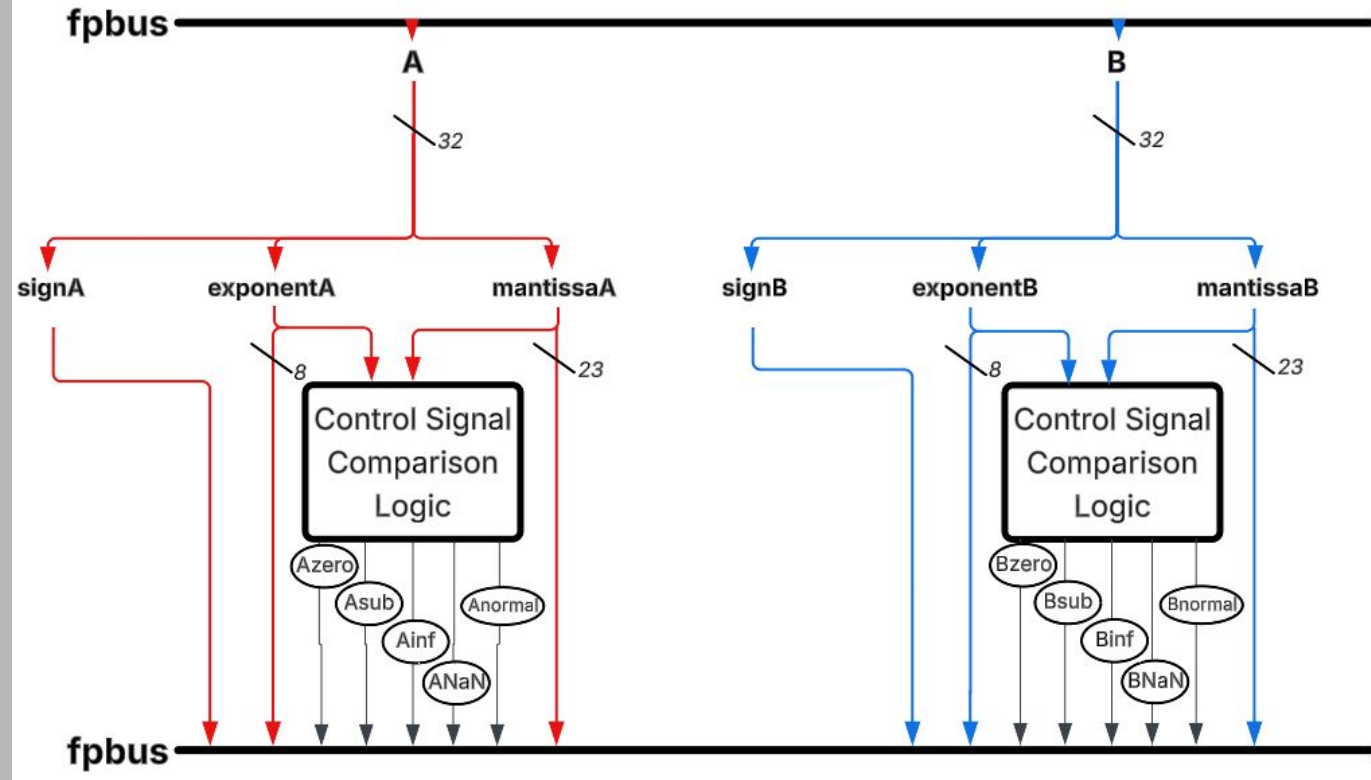
# Interface (*fpbus*)

- Common interconnect for all signals reused by multiple submodules
- Allows for built-in naming checks
- 33 total signals
- 5 modports (one for each submodule)



# Masking Module (Mask)

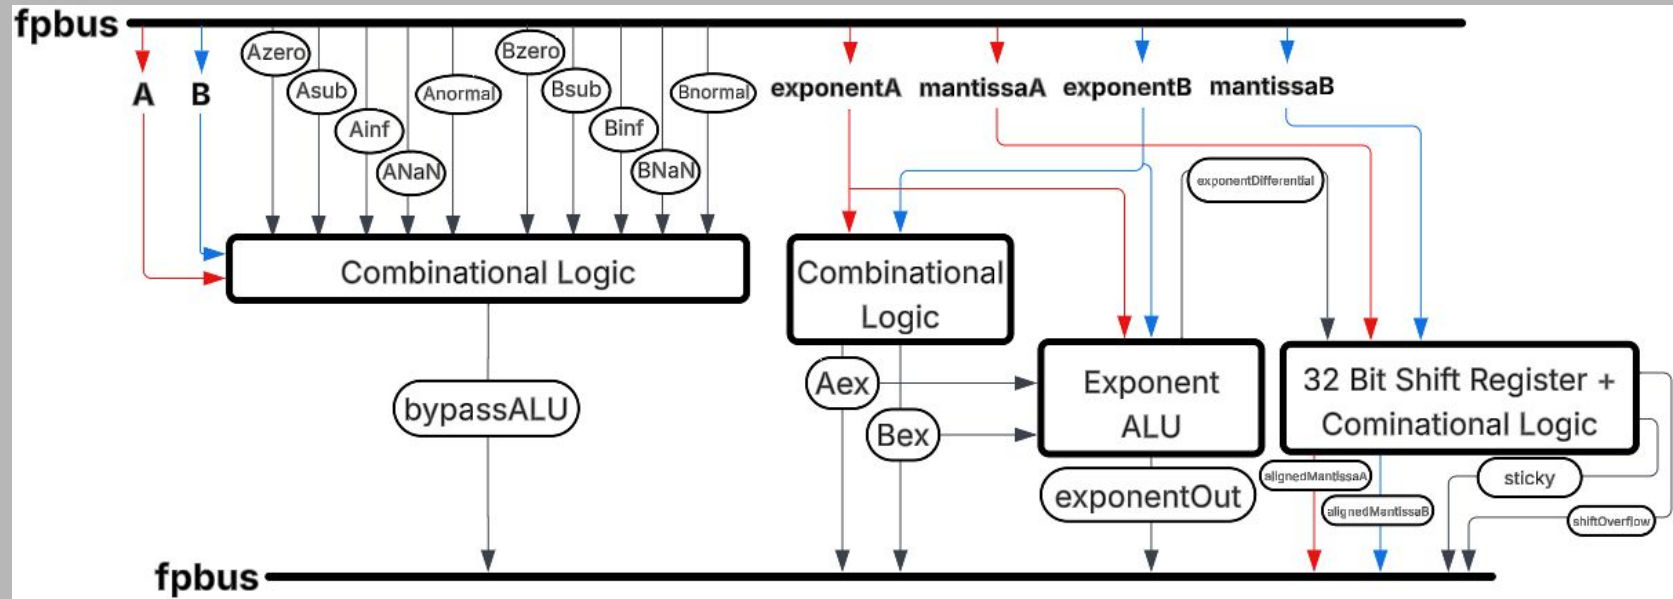
- Splits each input directly into a Sign, Exponent, and Mantissa value using bit slicing



- Computes and generates the control signals for both A and B to determine if there are non-normal numbers present in either addend

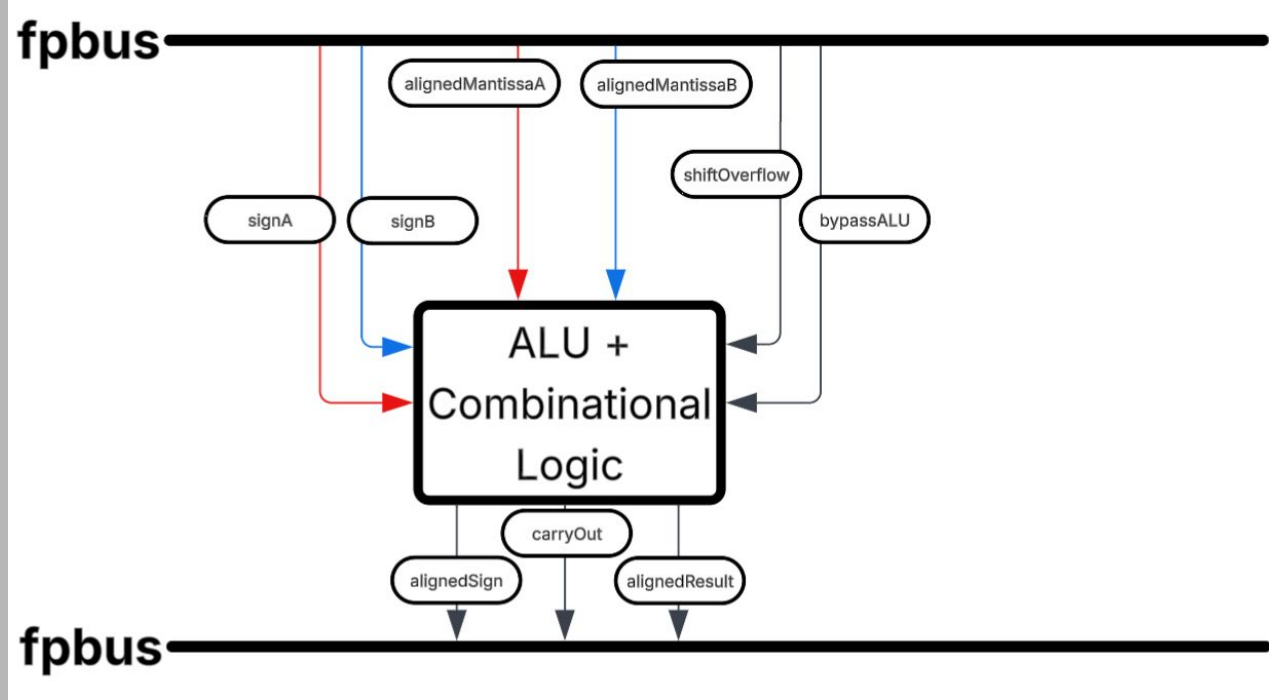
# Alignment Module (*Alignment*)

- Adds implicit ones and normalizes the lesser mantissa
- Signal for bypassing ALU for special cases such as NaN, Inf, Zeros
- Calculates Sticky Bit & Shift Overflow for proper rounding



# ALU Module (*ALU*)

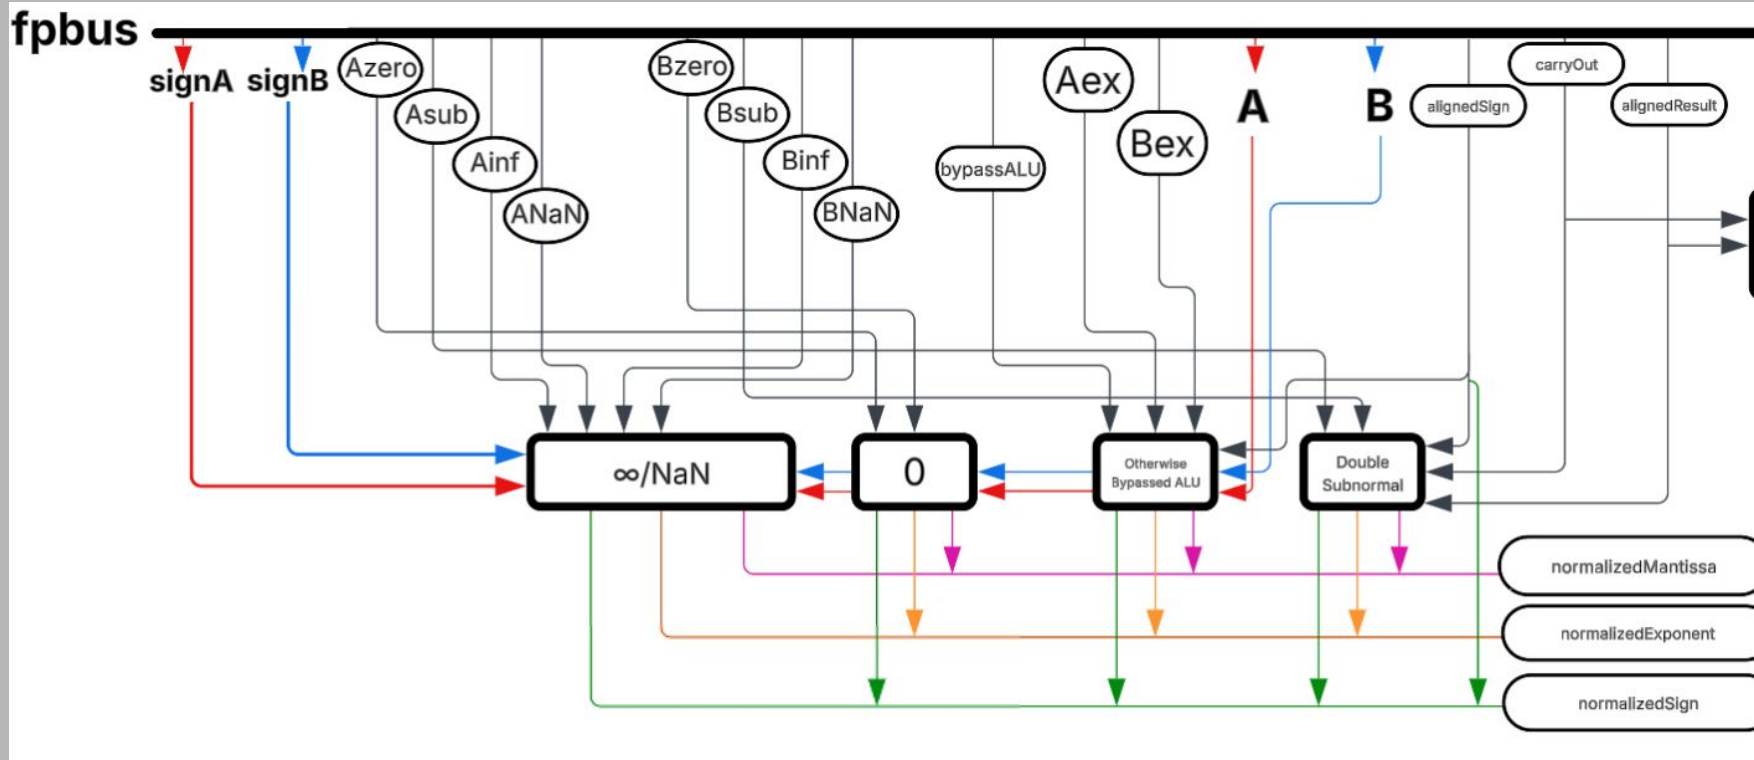
- Adds (or subtracts) the two aligned mantissas, plus additional subtraction for rounding accuracy



- Bypassed by special cases (handled during normalization)
- Tracks “carry out” for use in normalization
- Determines final sign bit to be used with a normal result

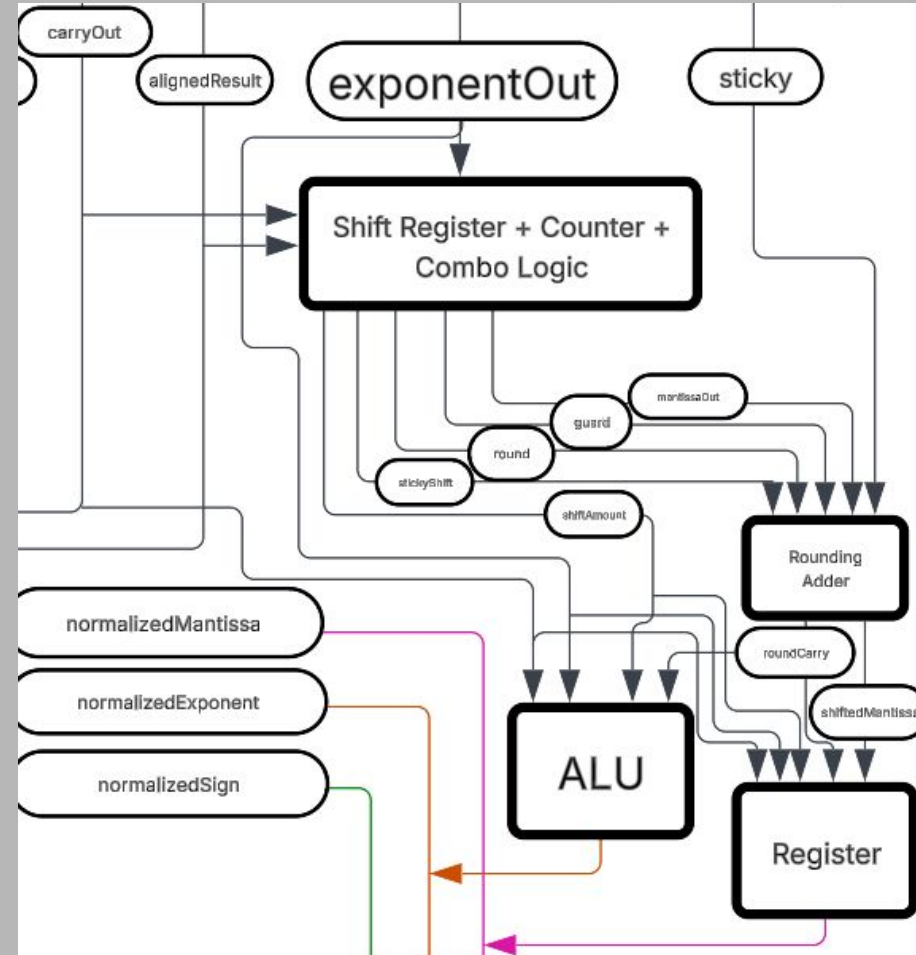
# Normalization Module (*Normalize*)

- Handles special cases that bypassed the ALU ( $\pm\infty$ ,  $\pm\text{NaN}$ ,  $\pm 0$ ,  $\pm\text{Subnormal}$ )
- If inputs result in normal number, this portion will be bypassed



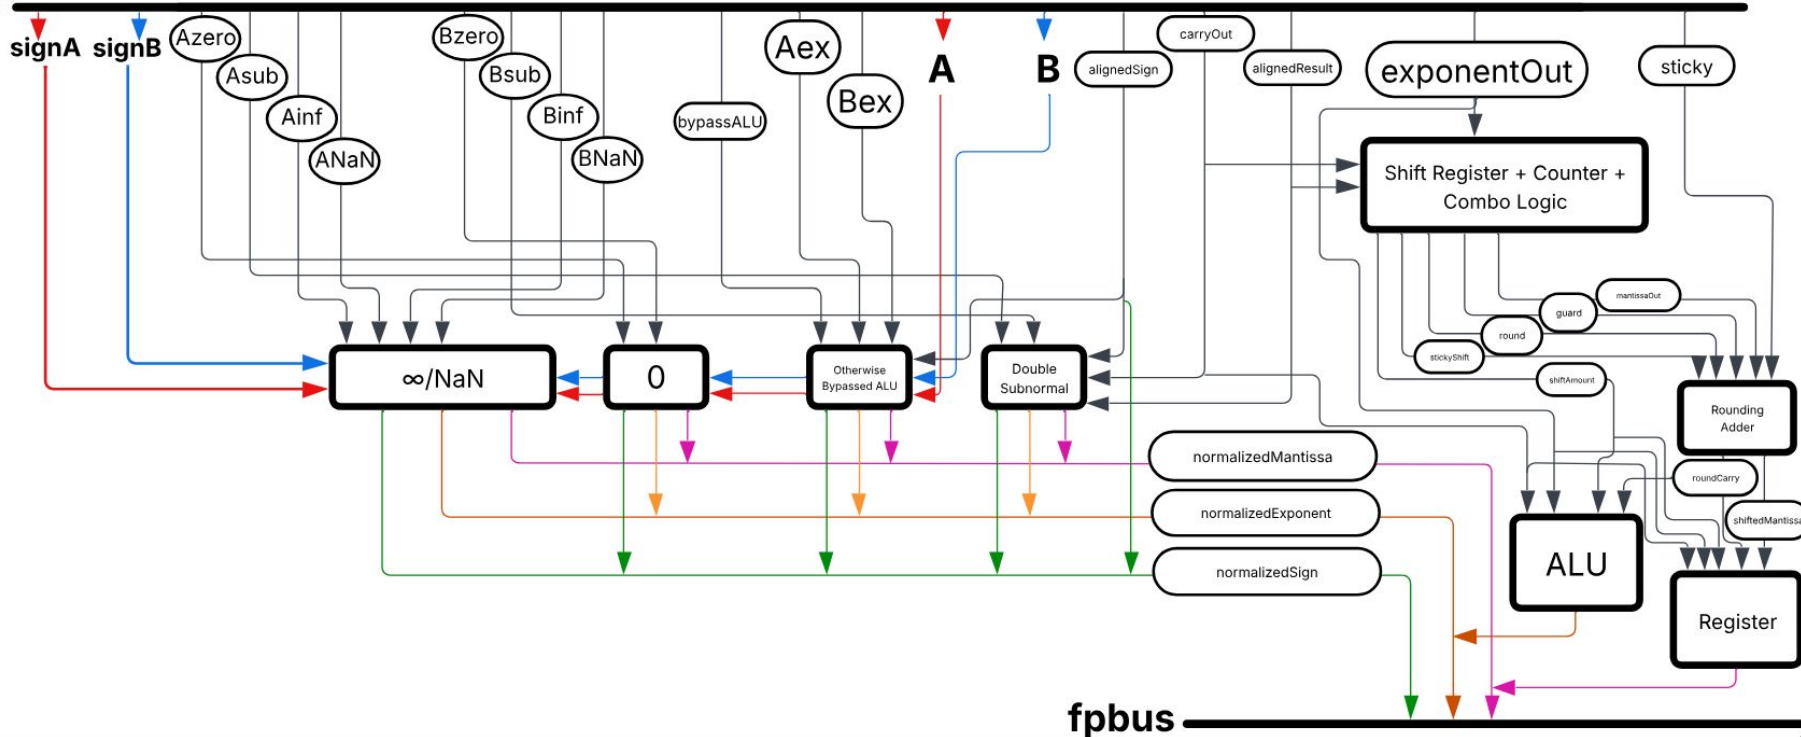
# Normalization Module (*Normalize*)

- Utilizes newly computed Sticky Bit, prior Sticky Bit, Round Bit, Guard Bit, and previous subtraction to accurately Round to Nearest “Even”
- Normalizes result by shifting and correcting Exponent
- Handles “carry outs” if mantissa shifts into the implicit one





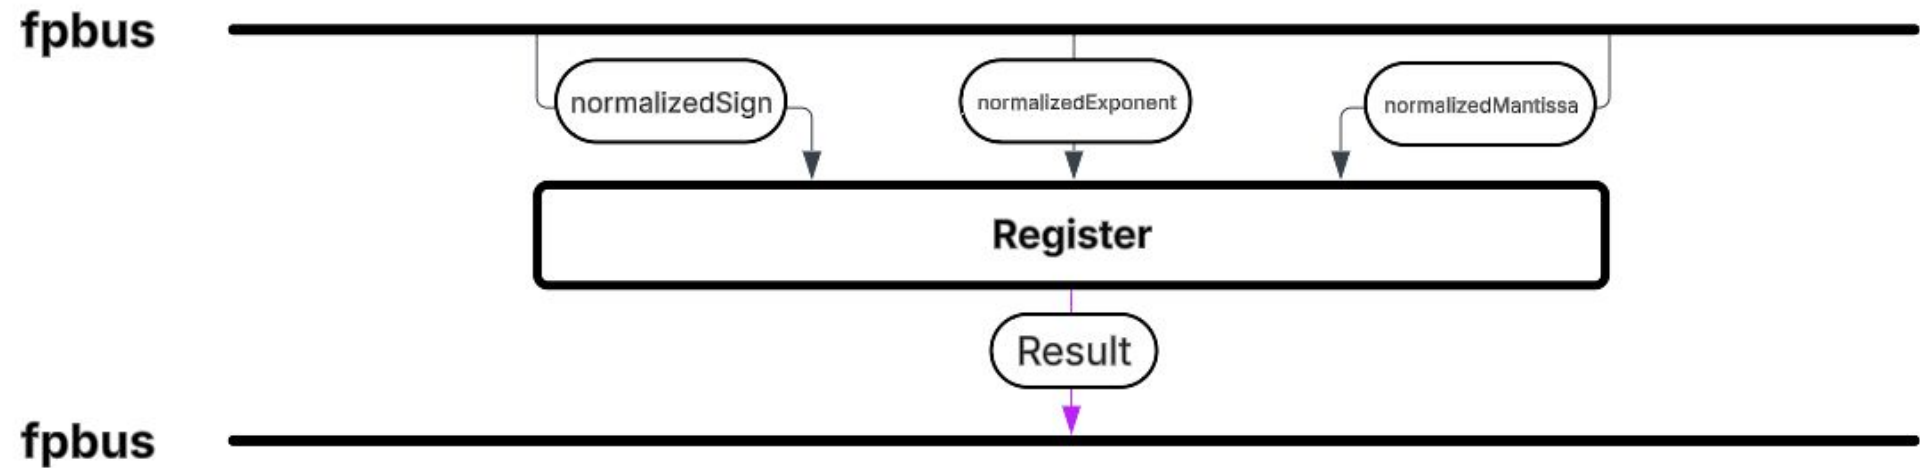
fpbus



fpbus

# Repacking Module (*Pack*)

- Very simple - concatenates normalizedSign, normalizedExponent, and normalizedMantissa into a single result for IEEE754 32-bit floating point representation



# Testing

- Every module was tested as we designed them to ensure no errors were made before proceeding (they still were!)

## *Mask-*

- Basic functionality test with  $2^{20}$  random 32-bit inputs

## *Alignment-*

- Uses *Mask* outputs as inputs, with  $2^{20}$  random test cases, including NaN, Infs, Subnormals, Zero and Normal Numbers to verify proper results and control signal assertion

# Testing

*ALU* -

- Determines expected exponent and mantissa separately, then checks those against *ALU* output (also computed using *Mask* and *Alignment*). Checks *bypassALU* signal to skip computation if necessary. Again,  $2^{20}$  test cases.

*Normalize* -

- One static test for early functionality, and  $2^{20}$  random tests, using all prior modules to compute all inputs to the *Normalize* module, and using *\$shortrealtobits* for results checking.

# Testing

## *Full System Testing-*

- 25 targeted tests cases (Case Analysis)
  - Done using a file, with hex value inputs for easy additions later on if deemed necessary
- $2^{25}$  constrained random non-NaN and non-Inf numbers
- $2^{25}$  random numbers including NaN and Inf
- Task for automatic results checking

# Bugs

## Minor Issues

- Dozens of syntax related bugs in early development (some during later development as well)
- Improper handling of non-normal numbers (by IEEE-754 conventions), which were fixed during *Normalize*

## Rounding Issues

1. We originally sampled GRS bits during the shifting process in *Alignment*, which resulted in rounding to ALWAYS be off by 1-3 (bits)
  - Moved to *Normalize*
2. Results still proved to be incorrect by 1-2 (bits)
  - Extended our ALU to use 32 bits, rather than the 26 we used previously (Mantissa + GRS), as a device capable of producing a 32-bit number almost definitely has 32-bit ALUs
  - Added second “Sticky Bit” still within alignment
3. Results were now consistently off by 1 bit
  - Added shift overflow check bit for 1’s lost with shifting, which accounted for any required subtraction that would usually be lost

# Results

- Accurate computation of “Normal” additions within precision
- Accurate rounding (adhering to IEEE-754 Round to Nearest “Even”)
- Handles special cases, Including NaN, +/- Inf, Subnormals (including those who graduate), etc
- Easily modifiable to support different rounding modes via loading different packages
- 0 Errors out of  $\sim 2^{26}$  test cases

What would our next additions be?

- Implementation of more rounding modes
- 64-bit & 16-bit modes, which would be toggled via a parameter

# Resources

[FP adder](#)

[FP converter](#)

[Floating Point Numbers and Rounding](#)

[IEEE 754 Floating Point Representation](#)

ECE 571 & 586 Lecture Slides